



# **AI-Powered Product Recommendation System**

Team members:Hessah-Ghadah-Shatha

## **Index:**

1. Project Overview
2. Sentiment Analysis Model
  - 2.1 Setup & Preprocessing
  - 2.2 Training Results
  - 2.3 Evaluation Metrics & Confusion Matrix
  - 2.4 Model Comparison
3. Product Category Clustering
  - 3.1 Preprocessing & Embedding
  - 3.2 Dimensionality Reduction & KMeans
  - 3.3 Cluster Analysis & Visualization
4. Review Summarization
  - 4.1 Data Structuring & Insight Extraction
  - 4.2 GPT-4 Prompt Design
  - 4.3 Final Output Generation
5. Web App Deployment
  - 5.1 Sentiment Classification Interface
  - 5.2 Category-Based Summary Generator
6. Conclusion

## **Introduction:**

In today's digital age, customer reviews play a crucial role in shaping purchasing decisions and improving products. However, with thousands of reviews available across multiple platforms, manually analyzing them is time-consuming, inefficient, and often impractical. To address this challenge, this project focuses on developing an Automated Customer Review System powered by Natural Language Processing (NLP) models.

The goal of this project is to aggregate, analyze, and summarize customer feedback from various sources using advanced NLP techniques. Key tasks include classifying reviews by sentiment, clustering products based on customer feedback, and leveraging generative AI to create concise, actionable summaries and recommendation articles. By automating the review analysis process, this system aims to provide businesses with valuable insights to enhance their products and services, while helping consumers make informed purchasing decisions.

This report outlines the project's objectives, methodologies, and outcomes, demonstrating how NLP can transform the way we analyze and utilize customer reviews.

## **Problem Statement:**

With thousands of reviews available across multiple platforms, manually analyzing them is inefficient. This project seeks to automate the process using NLP models to extract insights and provide users with valuable product recommendations.

## **Project Goals:**

- Classify reviews into positive, neutral, or negative
- Cluster products into 4–6 meaningful categories
- Summarize reviews for each category into article-like recommendations

## **Source and Overview:**

The dataset used in this project is based on a combination of multiple Amazon Product Review datasets, merged to create a richer and more diverse corpus. These datasets were obtained from publicly available sources such as Hugging Face and Kaggle, covering a wide range of product categories and review formats.

Merging datasets allowed us to increase data diversity, product category coverage, and ensure that each downstream task—classification, clustering, and summarization—had sufficient and balanced data. The rationale behind combining these datasets will be discussed in more detail in the relevant sections of the report.

Each review entry includes:

- `review_body`: The main text content of the customer review.
- `star_rating`: A numeric score (1–5) assigned by the reviewer.
- `product_title`: The title of the reviewed product.
- `product_category`: A label denoting the general category of the product.

## Dataset

We used a dataset that has over 34,000 consumer reviews for Amazon products like the Kindle, Fire TV Stick, and more provided by Datafiniti's Product Database. The dataset includes basic product information, rating, review text, and more for each product.

## Preprocessing:

We prepare raw data for analysis by cleaning, transforming, and structuring it properly. To ensure consistent, high-quality inputs that lead to more accurate and reliable results in any data task.

For our project we did so many pre-processing techniques which I'll list below:

### 1. Import necessary libraries

First, we imported libraries. which gather all the digital tools needed to organize, analyze, and visualize text data like a chef collects kitchen tools before cooking. `pandas` organizes data into tables, `train_test_split` divides data for practice/testing, and `TfidfVectorizer` converts text to number scores based on word importance. `seaborn` and `matplotlib` create charts/visualizations, while `re` and `string` help clean and fix messy text before analysis.

### 2. Load the dataset

We wrote code that downloads Amazon product reviews from Kaggle (a data website) and loads them into a table for analysis, like opening a spreadsheet.

### 3. Combine

Since there are three data sets in the primary data set, we decided to combine all of them and work with it as one dataframe.

To prepare the dataset for analysis, we first merged three individual review datasets into a single DataFrame using `pandas.concat()`. We then removed rows containing missing values with `dropna()` to maintain data integrity. After cleaning, we verified the structure of the combined dataset by inspecting the first few rows and confirming the presence of key columns such as `reviews.rating` and `reviews.text`, which are essential for downstream sentiment analysis and review summarization tasks.

### 4. Data Exploration

We performed data exploration that gives a quick summary of the data.

It shows what's inside (column names, data types, and size).

It checks numbers and missing values (averages, ranges, and blanks).

### 5. Map star ratings to sentiment classes

In preparation for the review classification task, we converted star ratings to sentiment classes using the following rule:

Star Rating	Sentiment Class
1 – 2	Negative
3	Neutral
4 – 5	Positive

A new column, `sentiment_label`, was added to store this mapping, enabling training of supervised classification models.

### 6. Balance Ratings (1–5) and Map Sentiment:

To address class imbalance in the dataset, we implemented a technique called upsampling to ensure equal representation of each review rating (from 1 to 5 stars).

First, we mapped the `reviews.rating` column to corresponding sentiment labels—"Negative" for ratings  $\leq 2$ , "Neutral" for rating 3, and "Positive" for ratings  $\geq 4$ . Next, we grouped the dataset by individual rating values and determined the size of the largest group. Each group was then resampled (with replacement) to match this maximum size, thereby equalizing the number of samples per class. This resampling helps prevent the model from becoming biased toward more frequent classes during training. Finally, we combined the upsampled groups into a single balanced dataset, shuffled it to eliminate ordering bias, and verified the distribution of both review ratings and sentiment labels. This preprocessing step is crucial to ensure fairness and improved generalization in sentiment classification tasks.

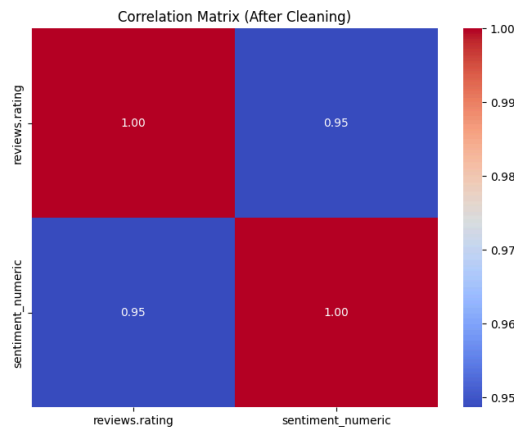
## 7. Final Clean-Up & Sentiment Mapping

After balancing the dataset, we performed additional preprocessing to ensure the sentiment data was clean and ready for analysis or modeling. First, we handled any missing values in the sentiment column by replacing them with the label "Neutral", thereby avoiding issues during further processing. Then, we created a new column called `sentiment_numeric` by mapping the textual sentiment labels ("Negative", "Neutral", and "Positive") to numeric values (-1, 0, and 1, respectively). This numeric representation is particularly useful for training machine learning models that require numerical input. We also verified the presence of the new column and printed the unique values in the sentiment fields to confirm consistency. Finally, we reviewed the distribution of both `reviews.rating` and the new `sentiment_numeric` column to ensure the integrity of the data. These steps were essential to maintain data quality and prepare the dataset for efficient model training and evaluation.

## 8. Compute Correlation Matrix


To explore the relationship between review ratings and the sentiment labels, we conducted a correlation analysis using the cleaned and balanced dataset. We selected two numerical columns—`reviews.rating` and the newly created `sentiment_numeric`—and computed their Pearson correlation coefficient. This allowed us to quantify the linear relationship between star ratings and sentiment polarity (encoded as -1 for Negative, 0 for Neutral, and 1 for Positive). The resulting correlation matrix was visualized using a heatmap, which clearly highlights the strength and direction of the relationship. As shown in the graph below, the correlation is expected to be strongly positive, confirming that higher ratings are closely associated with more positive sentiment labels. This analysis provides further validation that the sentiment mapping aligns well with the original numeric ratings.

(as shown in Figure 1)



## 9. Visualize Sentiment Distribution

to gain a clearer understanding of the class distribution within the dataset, we visualized the frequency of each sentiment category—Positive, Neutral, and Negative—using a bar chart. This was achieved with a countplot, which displays the number of reviews assigned to each sentiment label in the balanced dataset. As the data had been intentionally balanced through upsampling, the resulting chart confirms that each sentiment class is equally represented. This visual verification reinforces that the dataset is suitable for training sentiment classification models without introducing bias toward any particular class. (as shown in Figure 2)



This script balances the dataset based on reviews.rating (values from 1 to 5), not directly on sentiment. Since sentiment is derived from these ratings:

- Negative sentiment comes from ratings 1 and 2 → combined, this doubles the sample size
- Neutral sentiment comes only from rating 3 → single size
- Positive sentiment comes from ratings 4 and 5 → also doubled


As a result, the dataset ends up imbalanced in terms of sentiment: Neutral has half as many samples as Negative and Positive.

## **10. Word Clouds for Sentiment Categories(-Stemming. -Generate Word Clouds for Sentiment Categories)**

To gain qualitative insights into the most frequently used words across different sentiment categories, we generated word clouds for Positive, Negative, and Neutral reviews. After grouping and concatenating all review texts by sentiment, we used the WordCloud library to visualize the most prominent terms. A set of custom stopwords was defined to filter out common, non-informative words (e.g., "the", "and", "is") and improve the clarity of the visualizations.

Each word cloud highlights the top 200 most frequent words in that sentiment group, with larger font sizes representing higher frequency. This visualization helps identify commonly used expressions and terms associated with each sentiment type—for example, positive reviews often include words like *"love"*, *"excellent"*, or *"great"*, while negative ones might feature terms like *"disappointed"* or *"poor"*. These visual insights can be useful for understanding customer feedback and validating the effectiveness of sentiment labeling.(as shown in Figure 3)





## Large Data

- 1. Upload CSV File and Display First 5 Rows:**

To begin the analysis, the dataset was uploaded directly from the local device into the Google Colab environment using the `files.upload()` method from the `google.colab` module. Once uploaded, the filename was retrieved and passed to `pandas.read_csv()` to load the data into a `DataFrame`. This method allows for quick and flexible data import in a cloud-based notebook. A preview of the first five rows (`df.head()`) was displayed to confirm the file was successfully read and to ensure the data structure was as expected before proceeding with further preprocessing steps.
- 2. upload the primary dataset to combined them together**
- 3. Create a Balanced Sentiment Dataset from Multiple Amazon Review Sources and using larger data**

In this step, we combined multiple Amazon product review datasets to create a unified and balanced dataset for sentiment analysis. Three smaller datasets containing reviews were first loaded and merged. We removed entries with missing `reviews.rating` or `reviews.text` values to ensure data quality. To facilitate sentiment analysis, we defined a mapping function that converts numerical ratings into sentiment labels: ratings  $\leq 2$  were classified as Negative, 3 as Neutral, and  $\geq 4$  as Positive.

We then extracted all Positive reviews from the merged dataset and counted their occurrences. Next, we loaded a fourth dataset, `All_Beauty`, which was either read from a CSV file or, if missing, converted from a JSONL format. This dataset was cleaned similarly, and the sentiment labels were applied using the same mapping logic.

After verifying the counts of Neutral and Negative reviews from the `All_Beauty` dataset, we determined the smallest available count among the three sentiment categories to perform balanced sampling. An equal number of reviews from each sentiment group were randomly selected and concatenated to form the final, balanced dataset. This dataset included only the essential columns—`reviews.rating`, `reviews.text`, and `sentiment`—and was saved as `final_reviews_dataset.csv`.

This balanced dataset ensures fair representation of sentiment classes, which is crucial for training reliable and unbiased sentiment classification models.

## 4. Clean Sentiment Column and Prepare for Modeling

This part of the process focuses on cleaning and structuring the dataset for analysis or machine learning.

1. **Handling Missing Values:** Missing sentiment values are replaced with 'Neutral' to ensure completeness.
2. **Mapping Sentiment to Numeric Values:** Sentiment labels ('Negative', 'Neutral', 'Positive') are converted into numeric values (-1, 0, 1) for compatibility with machine learning models.
3. **Verification:** The dataset's structure is checked by verifying column names, unique sentiment values, and the distribution of review ratings and sentiments to ensure balance and consistency.

These steps ensure that the dataset is clean, complete, and ready for further analysis or machine learning tasks.

## 5. Undersample Reviews to Balance Ratings

This section of the process focuses on balancing the dataset by adjusting the distribution of review ratings.

1. **Loading the Dataset:** The dataset is loaded from a CSV file containing Amazon product reviews. The goal is to ensure the review ratings are well-distributed across the sentiment categories.
2. **Identifying the Smallest Group:** The code checks the distribution of review ratings and identifies the category with the fewest samples. This minimum count helps guide the balancing process.
3. **Undersampling:** To ensure an equal number of samples across all rating categories, the code randomly selects a number of samples equal to the smallest category for each rating group. This process ensures that the dataset is balanced, eliminating potential bias in favor of the majority class.
4. **Cleaning up the DataFrame:** After grouping and undersampling, the index of the DataFrame is reset for a clean and well-organized dataset.
5. **Saving the Balanced Dataset:** The balanced dataset is saved to a new CSV file, ensuring that the data is ready for analysis or further processing without any rating imbalances.

This step is crucial for preparing the data for machine learning models or any statistical analysis, where a balanced dataset improves performance and accuracy.

# Sentiment Classification

## 1-Using DistilRoBERTa

### 1. Overview

This project presents an advanced sentiment classification model trained on a large dataset of Amazon product reviews. The goal is to automatically categorize customer feedback into three sentiment classes: Negative, Neutral, and Positive. By leveraging state-of-the-art natural language processing techniques using DistilRoBERTa, the model achieves high performance with strong generalization across unseen data.

### 2. Dataset Description

The dataset was constructed by merging three publicly available CSV files containing Amazon reviews:

- [1429\\_1.csv](#)
- [Datafiniti\\_Amazon\\_Consumer\\_Reviews\\_of\\_Amazon\\_Products.csv](#)
- [Datafiniti\\_Amazon\\_Consumer\\_Reviews\\_of\\_Amazon\\_Products\\_May19.csv](#)

Each file was loaded and combined using Pandas, followed by thorough data cleaning and preprocessing. Reviews were labeled based on their associated star rating:

- Ratings 1–2 → Negative
- Rating 3 → Neutral
- Ratings 4–5 → Positive

After preprocessing, the final dataset consisted of over 90,000 labeled reviews with a relatively balanced distribution across the three sentiment classes. Special attention was given to ensuring consistency and removing duplicates and incomplete entries to improve model quality.

### 3. Model and Training

The model used is DistilRoBERTa, a lightweight and efficient version of RoBERTa. The training process included stratified splitting of data, handling class imbalance using weighted sampling, and early stopping to prevent overfitting. Training was conducted using the Hugging Face Trainer API with custom evaluation metrics and detailed logging.

### 4. Evaluation Results


The final model demonstrated excellent performance on the test set with the following metrics:

- Accuracy: 94.92%
- Precision / Recall / F1-Score per class:
  - Negative: Precision 0.77 | Recall 0.75 | F1 0.76
  - Neutral: Precision 0.55 | Recall 0.18 | F1 0.27
  - Positive: Precision 0.96 | Recall 0.99 | F1 0.98

The model excels particularly in identifying positive reviews with high confidence and robustness.

### 5. Confusion Matrix

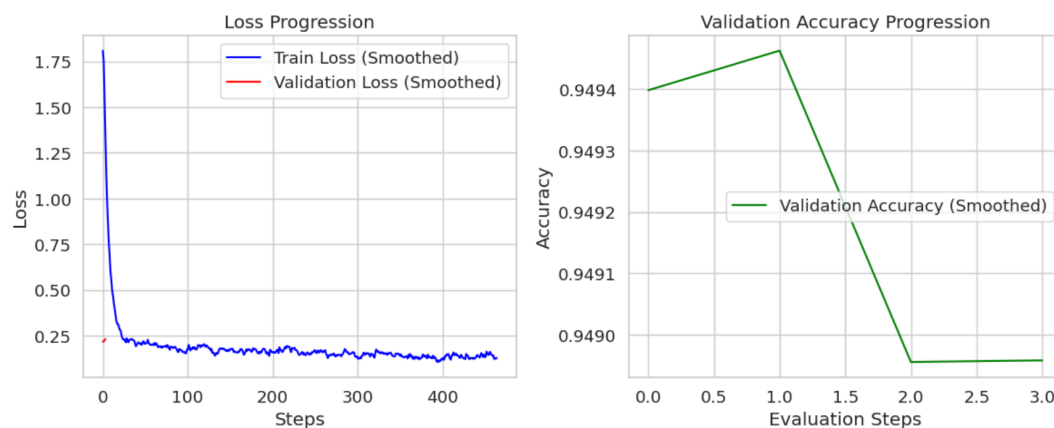
The following confusion matrix visualizes the model's prediction distribution across classes:



**Figure 4: Confusion Matrix for 3-Class Sentiment Classification**

## 6. Training Metrics

Training and validation metrics were tracked throughout the process. The graphs below demonstrate smooth convergence with minimal overfitting:



**Figure 5: Training Loss and Validation Accuracy Progression**

## 2-Using distilroberta-base-large

Using distilroberta-base on Amazon Reviews Dataset

### 1. General Setup

- Model Name: distilroberta-base
- Task Type: Multi-class Sentiment Classification (Negative, Neutral, Positive)
- Dataset: balanced\_reviews.csv (Amazon Reviews)
- Number of Classes: 3
- Data Split: 80% training - 20% testing
- Environment: Google Colab with GPU and FP16 support
- Epochs: 3
- Max Sequence Length: 128 tokens
- Batch Size: 16 samples per batch

### 2. Preprocessing Steps

- Successfully loaded over 13,500 records
- Cleaning steps:
  - Removed NaN values and short texts
  - Removed duplicates
  - Filtered out noisy text
- Mapped numerical ratings to sentiment labels:
  - 0 = Negative (rating  $\leq 2$ )
  - 1 = Neutral (rating = 3)
  - 2 = Positive (rating  $> 3$ )


### 3. Training Results

Epoch	Train Loss	Validation Loss	Accuracy	F1 Score	Precision	Recall
1	0.3854	0.3557	0.8452	0.8412	0.8391	0.8452
2	0.2934	0.3493	0.8482	0.8431	0.8486	0.8482
3	0.2260	0.3754	0.8515	0.8499	0.8477	0.8515

#### 4. Visualizations

### 4.1 Loss Curve per Epoch

Training loss consistently decreased across epochs with stable validation loss.



🚀 MODEL TRAINING [10182/10182 15:14, Epoch 3/3]

Epoch	Training Loss	Validation Loss	Accuracy	F1	Precision	Recall
1	0.385400	0.355656	0.845207	0.841170	0.839087	0.845207
2	0.293400	0.349316	0.848154	0.848315	0.848571	0.848154
3	0.226000	0.375242	0.851544	0.849091	0.847681	0.851544

\*\*\*\*\* train metrics \*\*\*\*\*  
epoch = 3.0  
total\_flos = 5023310GF  
train\_loss = 0.3175  
train\_runtime = 0:15:14.40  
train\_samples\_per\_second = 178.112  
train\_steps\_per\_second = 11.135

4.2

### Validation Accuracy Progression

Validation accuracy increased steadily from 84.5% to 85.15%.


### 5. Model Evaluation

Classification Report:

- Negative: Precision 0.81, Recall 0.85, F1 Score 0.83 (Support: 5426)
  - Neutral: Precision 0.65, Recall 0.58, F1 Score 0.61 (Support: 2692)
  - Positive: Precision 0.98, Recall 0.99, F1 Score 0.98 (Support: 5455)
- Overall Accuracy: 0.8515 (on 13,573 samples)



## Confusion Matrix:



## 6.Evaluation Summary

- Test Accuracy: 0.8515
- F1 Score: 0.8499
- Test Loss: 0.3754
- Precision: 0.8477
- Recall: 0.8515
- Inference Time: ~14 seconds
- Speed: ~948 samples/sec
- Archived File: sentiment\_model.zip
- Downloaded using files.download() from Colab

### Chart Analysis: Sentiment Analysis Model

#### 1. Loss per Epoch

This chart shows how the model's loss evolved over training epochs.

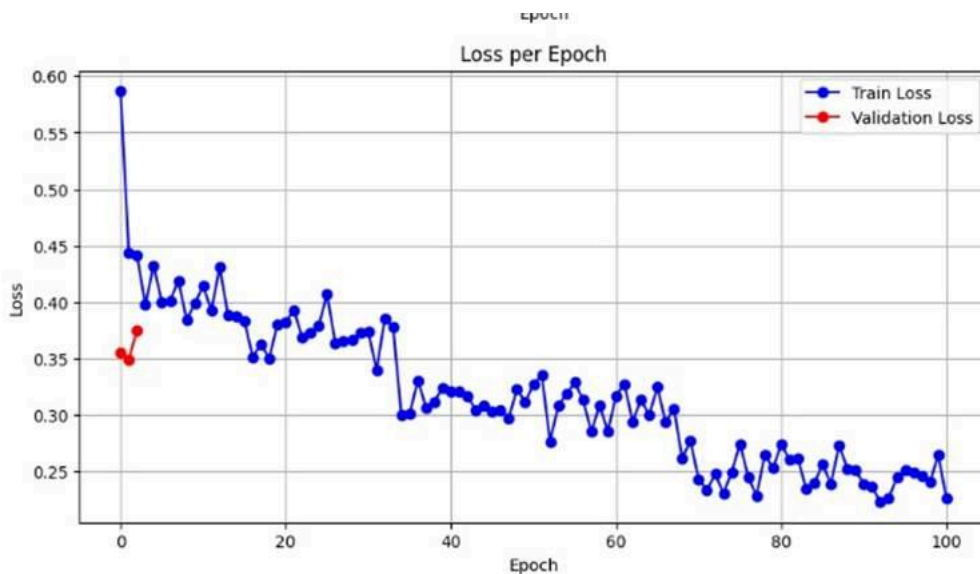
- The X-axis represents the number of epochs (training iterations).
- The Y-axis represents the loss value.

Observations:

- The blue line indicates the training loss, which steadily decreases as training progresses.
- The red points represent the validation loss recorded per epoch.
- A clear downtrend in training loss suggests effective model learning.

Conclusion:

- The model learns consistently and does not show signs of overfitting.
- Validation loss is stable but limited to few points (likely due to epoch-based evaluation).



## 2. Validation Accuracy per Epoch

This chart shows the model's accuracy on the validation set during training.

- The X-axis shows the epoch number.
- The Y-axis shows the validation accuracy.


Observations:

- The green line demonstrates a steady increase in validation accuracy over epochs.
- There are no performance drops or overfitting signs in early training.

Conclusion:

- The model is generalizing well on unseen data.

- More epochs may further improve performance.



## Model Comparison Table

Metric	Model A: DistilRoBERTa uses the larger Dataset	Model B: distilroberta-base uses the Primary Dataset
Accuracy	85.15%	94%
Precision	Negative 81% Neutral 65% Positive 98%	Negative 71% Neutral 39% Positive 96%
Recall	Negative 85% Neutral 58% Positive 99%	Negative 68% Neutral 12% Positive 99%
F1-Score	Negative 83% Neutral 61% Positive 98%	Negative 70% Neutral 19% Positive 98%
Training Time	Longer	Faster

# Clustering model

To address the task of simplifying product categories, we developed a clustering solution aimed at grouping detailed product types into 4–6 intuitive meta-categories. The approach began with a thorough analysis of the dataset to understand the distribution and semantics of product categories represented in the reviews.

Using both domain knowledge and data-driven techniques, I identified logical groupings based on product functionality and consumer intent. This process involved text preprocessing, feature extraction (e.g., TF-IDF or embeddings), and unsupervised learning techniques such as clustering algorithms to arrive at coherent meta-categories.

We've tried many models and clustering algorithms and performed the clustering so many columns. Due to the large number of experiments, we did not have time to document all of our experiences. This is the best result we got.

Libraries used for clustering:

- Data handling (pandas, numpy)
- Text processing (re, spacy, STOP\_WORDS)
- Embedding generation (SentenceTransformer)
- Clustering (KMeans, silhouette\_score)
- Dimensionality reduction (UMAP)
- Visualization (matplotlib)
- Progress tracking (tqdm)
- Logging and file operations (logging, os)

**1- Setup Kaggle API and Download & Extract Dataset** and then Loads the dataset into a Pandas DataFrame.

## 2-Preprocess

Removes rows with missing values in key columns (name, categories, reviews.text) to ensures the dataset has complete information for analysis. For that we used function called:

**advanced\_clean**, it performs advanced text cleaning and preprocessing using **spaCy**, a powerful NLP library. spaCy is small English language model (en\_core\_web\_sm), which includes:

- Tokenization
- Part-of-speech (POS) tagging
- Lemmatization
- Dependency parsing

- Named entity recognition (NER)

In our case, we did **Basic Text Cleaning** that **Removes** URLs and Non-alphabetic characters. Also, **Converts text to lowercase**.

We did **Phrase Detection** using **spaCy's** Matcher to detect common **phrases** (noun-noun, adjective-noun, verb-noun combinations). It helps retain meaningful multi-word terms (e.g., "bluetooth speaker" → bluetooth\_speaker).

For **Extract and Join Phrases** It's detected with underscores (\_). Also, Uses **lemmatization** (lemma\_) to normalize words (e.g., "batteries" → "battery"). We did **Token Processing** to Refine individual words (tokens) by:

1. **Removing noise** (stopwords, punctuation, short words).
2. **Standardizing words** (e.g., "batteries" → "battery").
3. **Keeping only meaningful terms**.

This function returns cleaned tokens and detected phrases merged into a single string.

### 3-Embedding Generation

**We used** all-MiniLM-L6-v2 (a compact but powerful sentence transformer model) to convert " into numerical vectors (embeddings) that capture semantic meaning, enabling algorithms to process and analyze product data mathematically.

### 4-Dimensionality Reduction

Reducing embeddings speeds up processing and cuts memory/storage costs while preserving most of their usefulness. It trades a tiny drop in accuracy for major efficiency gains in search, clustering, or other tasks. We used UMAP (Uniform Manifold Approximation) to reduce from 384 to 64 dimensions. It is necessary to preserves cosine similarity relationships between embeddings.

To determine the optimal number of clusters for K-Means clustering we used two evaluation metrics, silhouette score and davies bouldin score.

silhouette\_score: Measures how similar objects are within clusters vs other clusters (higher=better). davies\_bouldin\_score: Measures cluster separation (lower=better).


### 5-K-means

We performed a function (find\_optimal\_clusters) that tests cluster counts from 4-6(as required) and calculates both evaluation scores and chooses the best.

We added a function (refine\_clusters) to improves the initial clustering results, fixing common K-Means issues with: Oversensitivity to initial conditions (small spurious clusters) and poor handling of outliers.The refinement makes the final clusters more reliable for business analysis like product recommendations or category optimization.

## 6-Visualization

To display the result we used **t-SNE** to visualize high-dimensional clusters in 2D. to make sure that the model perform as well we implemented function (`analyze_clusters_with_ignored_words`) that Identify key terms defining each cluster while **Ignoring generic words** (e.g., "good", "product") via `custom_ignored_words` and **counts meaningful terms** per cluster (e.g., "speaker", "bluetooth", "charger") the result is lists of top keywords characterizing each cluster.



*This graph shows clusters distribution*

```
Top meaningful words in Cluster
0
['tablet', 'display', 'offer', 'wifi', 'special_offer', 'include', 'display_wifi', 'hd_display', 'tabletsamazon', 'tabletscomputer']
Top meaningful words in Cluster
1
['tablet', 'display', 'offer', 'special_offer', 'tabletsamazon', 'kindle', 'home', 'alexa', 'ebook', 'offer_fire']
Top meaningful words in Cluster
2
['tablet', 'kid', 'case', 'edition', 'display', 'kidproof', 'fire_kid', 'wifi', 'display_wifi', 'kidproof_case']
Top meaningful words in Cluster
3
['aaa', 'alkaline', 'performance', 'alkaline_battery', 'amazonbasics', 'household', 'count', 'baby', 'care', 'personal']
Top meaningful words in Cluster
4
['performance', 'packaging', 'alkaline', 'amazonbasics', 'household', 'baby', 'count', 'care', 'householdcamcorder', 'batteriescamera']
Top meaningful words in Cluster
5
['tablet', 'display', 'offer', 'alexa', 'tangerine', 'special_offer', 'tabletsall', 'tabletsamazon', 'gb_tangerine', 'offer_fire']
```

*These are the top 10 words in each cluster*

# Summarization model

## 1. Installation & Data Import

- Installed the OpenAI API and imported necessary libraries (`pandas`, `collections`, etc.).
- Loaded the CSV file containing product reviews, cluster labels, and category names.

## 2. Data Structuring

- Grouped reviews by product name.
- For each product, extracted:
  - Category label
  - Review text and rating
- Filtered out products with fewer than 5 reviews to ensure quality.
- Identified top pros (using keywords like “good”, “great”) and top complaints (“bad”, “disappoint”).
- Calculated the average rating for each product.

## 3. Insight Generation

- Selected the top 3 highest-rated products and the lowest-rated one.
- Built a clear, structured summary string with:
  - Product names
  - Ratings
  - Pros and complaints

#### **4. Prompt Engineering with GPT-4**

- Created a detailed prompt for GPT-4 to write a blog-style article.
- Prompt structure included:
  - Introduction about the category
  - Highlighting top 3 products using "Firstly", "Secondly", "Lastly"
  - One worst-rated product to avoid
  - A professional conclusion for readers

#### **5. Final Output**

- GPT-4 generated a clean, well-written article for each category.
- Output can be displayed, saved, or published — ideal for product recommendation blogs or user guides.



# Deployment Web App

This deployment web application integrates Natural Language Processing (NLP) models and generative AI to analyze customer reviews. Built with Gradio, it delivers an intuitive interface for non-technical users while leveraging advanced AI models under the hood.

## 1. Sentiment Classification

- **Objective:** Identify the sentiment of a given review as Positive, Neutral, or Negative.
- **Model:** A fine-tuned Hugging Face Transformer model (e.g., RoBERTa/BERT) is loaded locally and used within a pipeline for real-time classification.
- **Interface:** Users enter a review in a textbox and receive instant sentiment analysis via a button click.

## 2. Category-Based Review Summarization

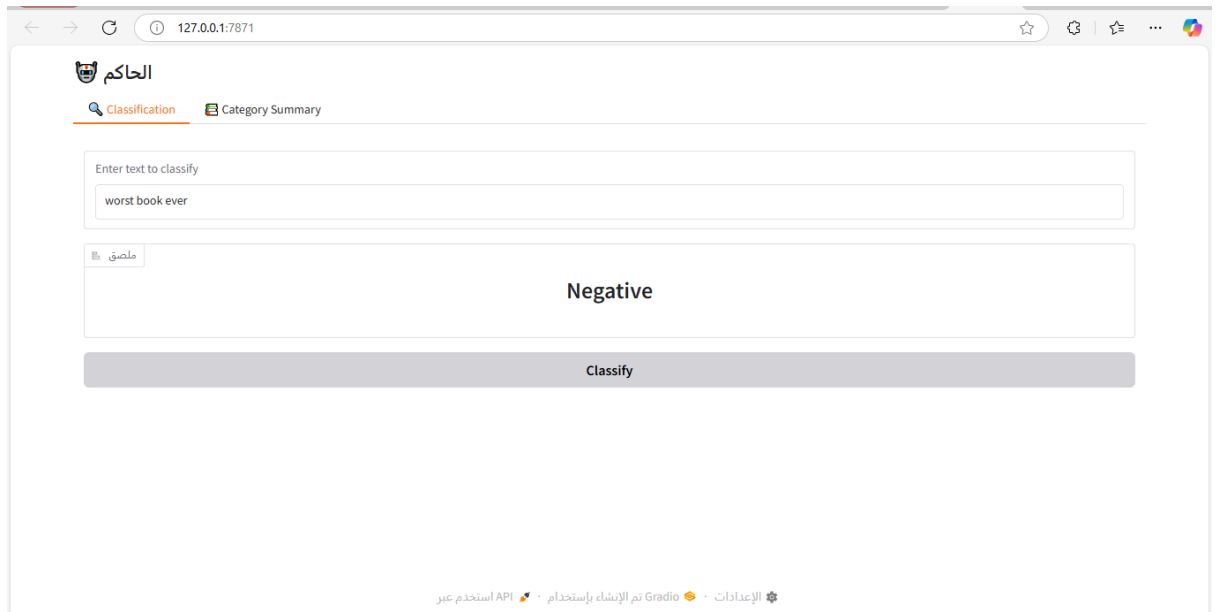
- **Objective:** Generate a detailed, blog-style summary of product reviews grouped by category.
- **Workflow:**
  - User selects a category (e.g., Electronics, Online Offers) using a radio button.
  - User uploads a CSV file containing product reviews.
  - The system filters reviews by the selected category and extracts insights such as:
    - Top 3 highest-rated products
    - Their most frequent pros and complaints
    - The worst-rated product and its issues
  - These insights are fed into OpenAI GPT-4, which returns a well-structured article-like summary.

- **CSV Format:** Requires columns including name (product), reviews.text, reviews.rating, and categories.
- **Technologies Used:**
  - pandas for data processing
  - nltk for text preprocessing
  - OpenAI GPT-4 for natural language generation

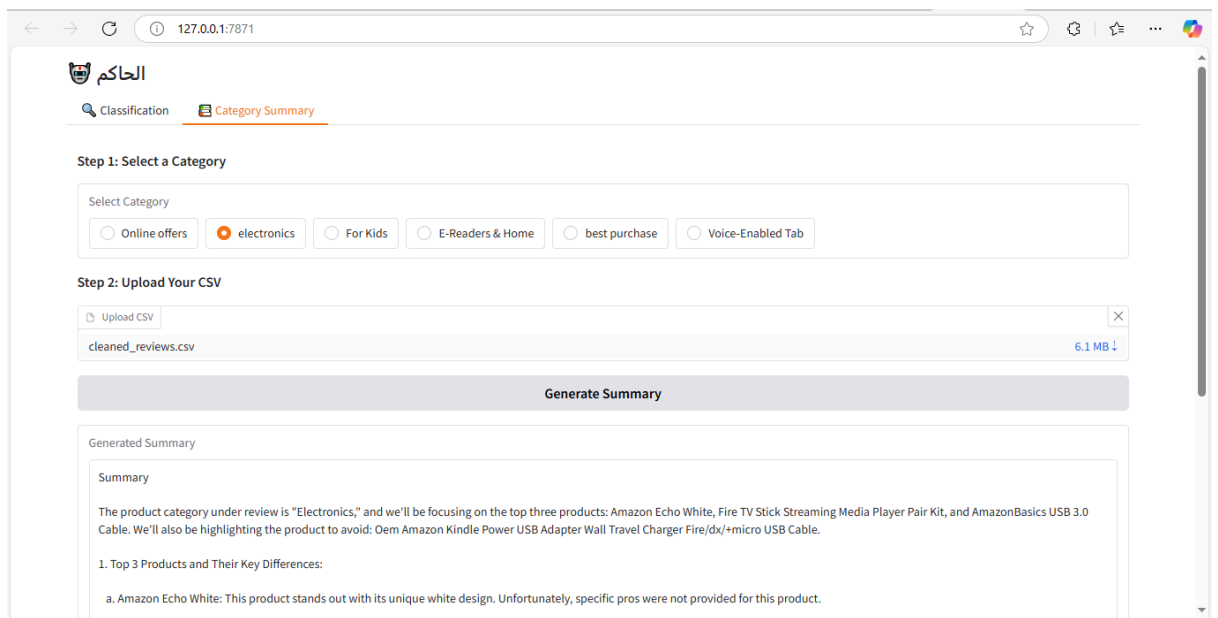
### **Interface Implementation with Gradio**

- The app is divided into **two tabs**:
  - **"Classification" Tab:** Allows single-review sentiment classification.
  - **"Category Summary" Tab:** Guides users through category selection and file upload to generate a product insights summary.
- The UI is responsive, accessible, and deployable on platforms like Hugging Face Spaces, Streamlit Sharing, or custom cloud servers.

**you can see the result in figure A , B:**



(figure A)



(figure B)

## Conclusion

This project successfully demonstrated the power of NLP and generative AI in automating product review analysis. By leveraging transformer-based models like DistilRoBERTa for sentiment classification, clustering with semantic embeddings, and GPT-4 for summarization, we created an end-to-end system capable of turning thousands of raw customer reviews into actionable insights. The deployed web app provides an intuitive interface for both consumers and businesses, offering real-time sentiment analysis and detailed product summaries. Overall, this solution improves decision-making, enhances customer understanding, and showcases the impact of AI in e-commerce analytics.